

# Entanglement-Efficient Bipartite-Distributed Quantum Computing

Prepared by Gao Dingchao

Reading Seminar

May 24, 2025

# References



Pablo Andres-Martinez, Senior R&D Scientist at Quantinuum



J-Y. Wu *et al.*, *Entanglement-efficient bipartite-distributed quantum computing*, Quantum **7** (2023).

# Outline

## 1 Background

## 2 Problem in Telegate

- Qubit Allocation
- Gate Distribution

## 3 WIP

# Why Distributed Quantum Computing (DQC)?

- **NISQ limitation:** single QPU constrained by qubit number, coherence, connectivity.
- **Connect QPUs via entanglement**  $\Rightarrow$  larger logical device.
- **Bottleneck:** entanglement distribution is costly & probabilistic.
- Goal:

Minimise *EPR pairs* consumption for two-party (bipartite) DQC while retaining universality.

# basic protocol







Figure: basic protocol for teleporting a qubit

$$|\psi\rangle |\beta_{00}\rangle \rightarrow \cdots \rightarrow \frac{1}{2} [ |00\rangle |\psi\rangle + |01\rangle X |\psi\rangle + |10\rangle Z |\psi\rangle + |11\rangle XZ |\psi\rangle ]$$

# Protocol in Distributed Quantum Operations



**Figure: Teledata protocol.** The quantum state is first teleported, and then the operation  $U$  is applied on the remote system.



**Figure: Telegate protocol.** The remote party applies a gate  $U$  using classical control based on measurement outcomes, without teleporting the quantum state.

# Newest Distributed Qubit Experiments (as of May 2025)

Platform	Teledata	Telegate
<b>Superconducting</b>	64 m cryo-bus state teleportation, Qiu <i>et al.</i> , 2025	99% SWAP/CZ via detachable cable, Mollenhauer <i>et al.</i> , 2025
<b>Trapped Ions</b>	14 km urban-fibre teleportation, Wang <i>et al.</i> , 2024	Teleported CZ & Grover's search (2 m), Main <i>et al.</i> , 2025
<b>Neutral Atoms</b>	420 km atomic-ensemble entanglement, Luo <i>et al.</i> , 2025	Not yet demonstrated

# Outline

## 1 Background

## 2 Problem in Telegate

- Qubit Allocation
- Gate Distribution

## 3 WIP



# The DQC Problem: Formal Definition

**Input:** Circuit  $C$  on qubits  $Q$ , architecture graph  $G = (V, E)$ .

- Each module  $A \in V$  has capacity  $\omega(A)$  (data qubits) and  $\varepsilon(A)$  (link qubits).
- Non-local gates (across modules) require 1 ebit each.

**Output:** A distribution  $(\varphi, \tilde{C})$  such that:

- Qubit allocation map  $\varphi : Q \rightarrow V$ ,  $|\varphi^{-1}(A)| \leq \omega(A)$ .
- $\tilde{C}$  includes EPR pairs needs for each non-local gate.
- Active Pairs  $\leq \varepsilon(A)$  at all times.

**Goal:** Minimise total ebit usage.

# Two Key Subproblems in DQC

- **Qubit Allocation:**

- Partition qubits across modules respecting  $\omega(A)$ .
- Minimise cut edges (i.e., potential non-local interactions).

- **Non-local Gate Distribution:**

- find a way to implement the non-local gates

# TC vs. QCD: Key Differences

- **Gate Types:** TC limits operations to those on-chip; QCD allows all, but nonlocal ones are costly.
- **Objective:** TC minimizes depth; QCD minimizes cross-QPU communication.
- **Optimization:** TC is local (gate-by-gate); QCD is global (qubit grouping).
- **Result:** TC outputs topology-compliant circuits; QCD allows nonlocal gates when needed.

# Formal Restatement as Partitioning

## Interaction Graph:

- Vertices: qubits  $Q$ .
- Edges: between  $q_1, q_2$  if a two-qubit gate in  $C$  involves them.

**Objective:** Partition  $Q$  into  $|V|$  disjoint sets

- Each set has size at most  $\omega(A)$ ,
- Minimise the number of cut edges (gates across different sets).

This is a **capacity-constrained graph partitioning** problem.

# Main Partitioning Strategies

- **Local Methods:**

- Iteratively improve an initial partition.
- Sensitive to starting configuration.
- Examples: Kernighan–Lin, Fiduccia–Mattheyses.

- **Global Methods:**

- Use global graph properties to guide partitioning.
- Avoid arbitrary initialisation.
- Examples: Spectral Partitioning, Multilevel Partitioning.

- **Simulated Annealing, ...**

# Considerations in Practice

- Partitioning quality affects EPR cost and circuit depth.
- Hardware-specific capacity limits must be respected.
- Gate commutativity can allow gate packet fusion before partitioning.

# Definition: Hypergraph Partitioning

**Definition:** A hypergraph  $G = (V, H)$  consists of a set of weighted vertices  $V$  and a set of hyperedges  $H$ , where a hyperedge  $h \subseteq V$  may connect more than 2 vertices.

**Connectivity metric  $(\lambda - 1)$ :**

$$\lambda_G = \sum_{h \in H} ((\# \text{partitions containing endpoints of } h) - 1)$$

The hypergraph partitioning problem is to find a balanced  $k$ -way partition minimizing this metric.

# Hypergraph Example




Figure: Example quantum circuit, applying Hadamard gates H and CZ gates.




Figure: hypergraph of the example circuit



# DQC and Hypergraph Mapping

## Correspondence between DQC and hypergraph partitioning:

Hypergraph	DQC
Vertex	Qubits $\cup$ CRZ gates
Hyperedge	Wire segment $\{q_i\} \cup \{g_1, \dots\}$
Vertex weight	1 for qubits, 0 for CRZ
$(\lambda - 1)$ metric	EPR cost
Partitioning	Qubit allocation and gate execution

# Why Hypergraph, Not Graph?

- Graph cut counts *each gate* separately.
- Hypergraph cut counts *each packet* once.
- Captures the cost structure of shared EJPPs (ebit sharing).




Figure: An example circuit




Figure: **Graph:** This partition shown cuts three edges




Figure: **Hypergraph:** This partition shown cuts One edges

# Key Result

## Theorem

For fully connected networks,

- Each valid distribution  $\Rightarrow$  a hypergraph partition
- If the circuit's hypergraph can be partitioned cutting  $n$  edges, the circuit can be distributed using  $n$  ebits.
- *Optimal partition  $\Leftrightarrow$  Minimum ebit allocation*

# Qubit Allocation Workflow Summary

- 1 Rebase circuit to  $\{H, RZ, CRZ\}$
- 2 Construct hypergraph (qubit/gate vertices, one edge per packet)
- 3 Use hypergraph partitioner (e.g. KaHyPar) with capacity constraints
- 4 Read module assignment from partition

**Ebit count = Hypergraph cut cost**

# Recall protocol

Controlled-unitary  $CU$  can be implemented non-locally with **one ebit** using:

- *Starting process*: cat-entangler.
- *Kernel*: local  $C_{e,U}$ .
- *Ending process*: cat-disentangler.




Figure: Telgate protocol

# Setting the Stage

Goal of gate distribution: realise every non-local gate while *minimising entanglement (ebit) cost*.




Figure: Examples of EJPP Protocol

# Distributing and Embedding Processes

## Definition: Distributing Process

A unitary  $U$  is *q-rooted distributable* if it is diagonal/antidiagonal in  $q$  and decomposes as  $U = \sum_{ij} \Delta_{ij} |i\rangle\langle j|_q \otimes V_j \otimes W_j$ .

- Implementable with **one ebit** (extends EJPP).

## Definition: Embedding Process

A unitary  $U$  is *q-rooted embeddable* if  $C_{q,X_e} U C_{q,X_e} = (L_A \otimes L_B) U (K_A \otimes K_B)$ .

- Enables merging of non-sequential distributing gates.

# Merging Packing Processes

**Statement:** If  $P_{q,e}[K_1]$  and  $P_{q,e}[K_2]$  are packing processes implementing  $U_1$  and  $U_2$  respectively, then

$$U_2 U_1 = P_{q,e}[K_2 K_1] \quad (\text{still 1 ebit}).$$



# Merging Packing Processes

**Statement:** If  $P_{q,e}[K_1]$  and  $P_{q,e}[K_2]$  are packing processes implementing  $U_1$  and  $U_2$  respectively, then

$$U_2 U_1 = P_{q,e}[K_2 K_1] \quad (\text{still 1 ebit}).$$

## Implication

Sequential or *merged-via-embedding* distributable gates can share **one** entangled pair.

# Conflict Graph and Edge Types

- **Vertices:** indecomposable packing kernels ('D' for distribute, 'B' for embed).
- **Edges:**
  - DD-type intrinsic conflict between two distributing options.
  - DB-type extrinsic resource conflict distribute vs embed.
  - BB-type embed vs embed (resource or structural).




Figure: DD-type conflict example.

# Steiner-Tree Entanglement Routing

- Prepare Bell pairs only along the edges of the *minimum Steiner tree* connecting the two modules.
- Keep the tree's entanglement *alive* for the whole packet  $\Rightarrow$  edges are created once, then reused.

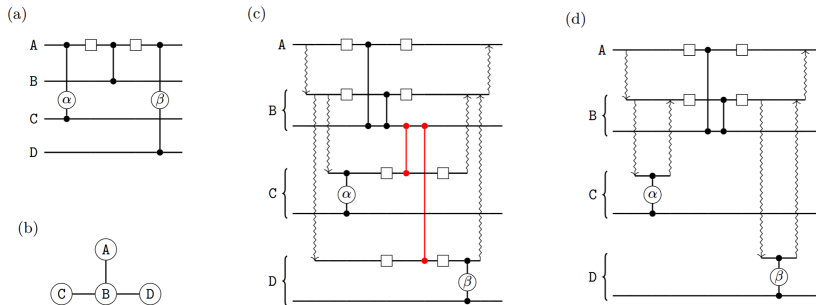


Figure: Steiner tree Example

# End-to-End Workflow


- ① **Allocate qubits** via hypergraph partitioning.
- ② **Group non-local CRZs** on each root into packets.
- ③ For each packet:
  - ① Find Steiner tree between the two modules.
  - ② Execute packet using **EJPP** over that tree.
- ④ Apply **embedding** and **detached** rewrites when profitable.

**Result:** full distributed circuit with (near-)minimal ebit count.

# Outline

- 1 Background
- 2 Problem in Telegate
  - Qubit Allocation
  - Gate Distribution
- 3 WIP

# Quantum Fourier Transform (4-Qubit Example)



**Figure:** Example 4-qubit QFT circuit. Controlled-phase gates  $P_\theta$  apply  $|1\rangle\langle 1| \otimes e^{i\theta Z}$ .

# Static $k$ -Partition Scheme





Figure: Illustration of a 3-partition mapping of the 4-qubit QFT.

# With Mid-Circuit Measurements



**Figure:** Dynamic variant: measurement outcomes on partitions A and B classically steer subsequent inter-partition operations.